1. **In this question (1), we will use attached dataset as nba.csv.**

Find the output of the following sections:

i)Write a syntax to load given data

Answer:

```
import pandas as pd
data = pd.read_csv("nba.csv")
data.head()
```

OUTPUT

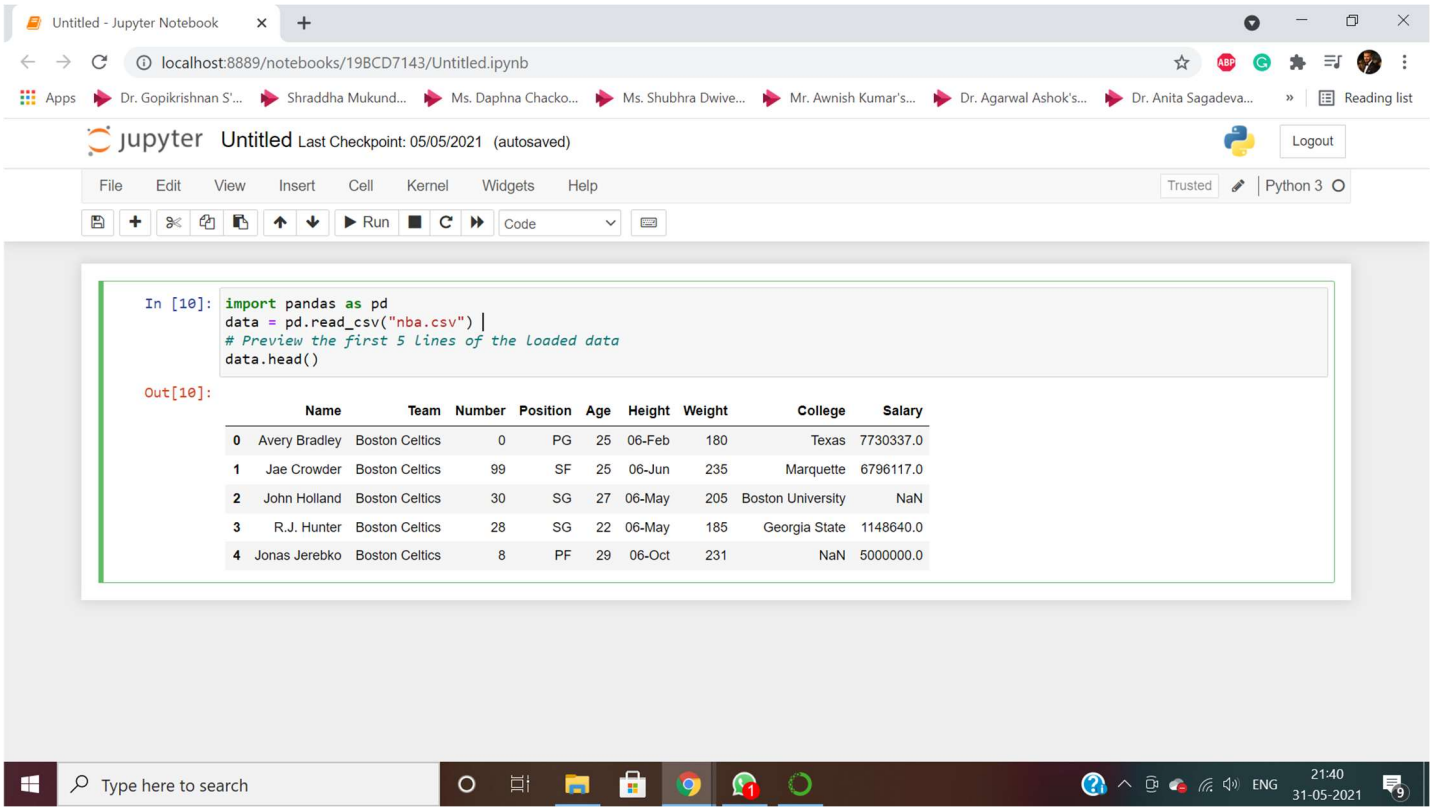

	Name	Team	Number	Position	Age	Height	Weight	College	Salary
0	Avery Bradley	Boston Celtics	0	PG	25	06-Feb	180	Texas	7730337.0
1	Jae Crowder	Boston Celtics	99	SF	25	06-Jun	235	Marquette	6796117.0
2	John Holland	Boston Celtics	30	SG	27	06-May	205	Boston University	NaN
3	R.J. Hunter	Boston Celtics	28	SG	22	06-May	185	Georgia State	1148640.0
4	Jonas Jerebko	Boston Celtics	8	PF	29	06-Oct	231	NaN	5000000.0

ii)Summarize the data

Answer:

```
import pandas as pd
data = pd.read_csv("nba.csv")
data.describe()
```

OUTPUT

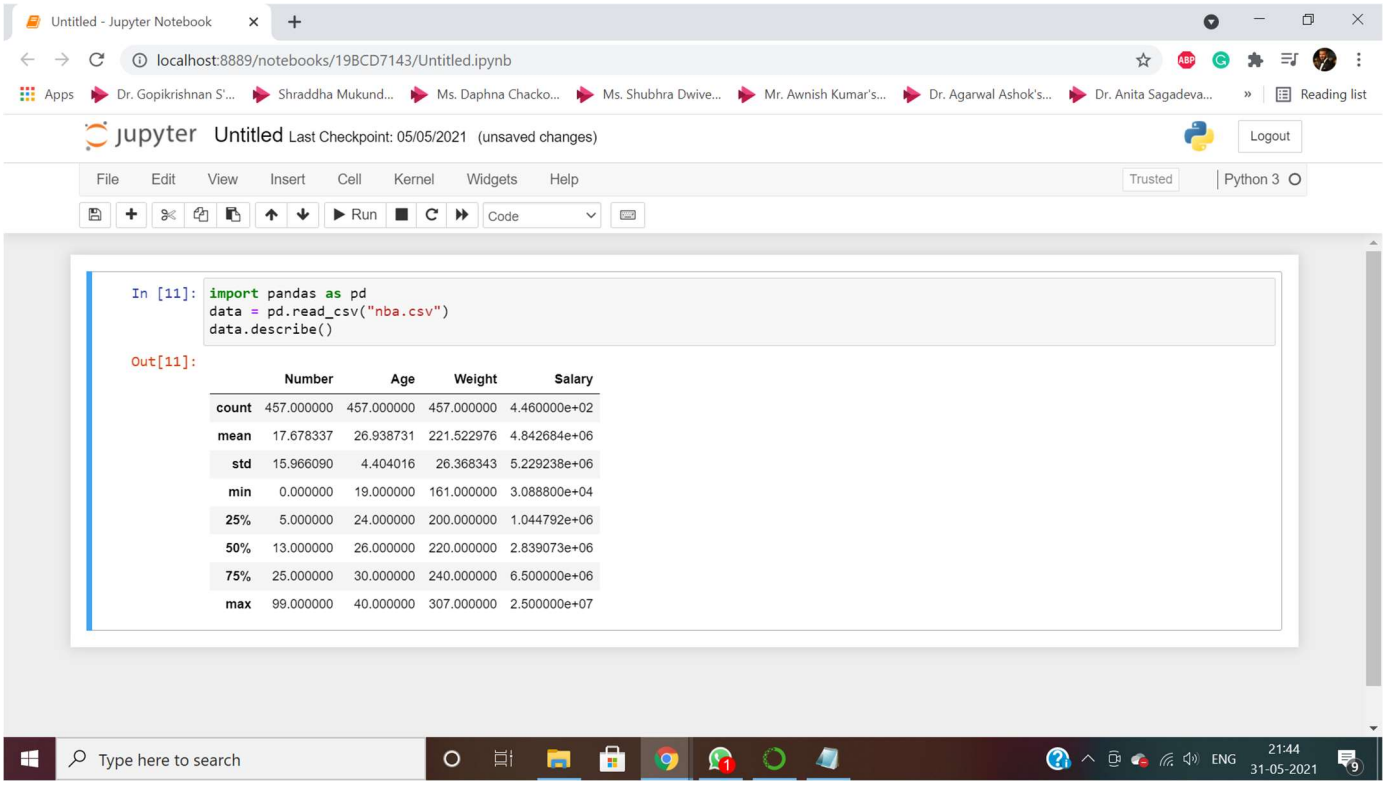

iii) To select columns Age, College and Salary for only rows with a labels Amir Johnson and Terry Rozier.

Answer:

```
import pandas as pd
data = pd.read_csv("nba.csv", index_col ="Name")
first = data.loc[["Amir Johnson", "Terry Rozier"],
                 ["Age", "College", "Salary"]]
first
```

OUTPUT

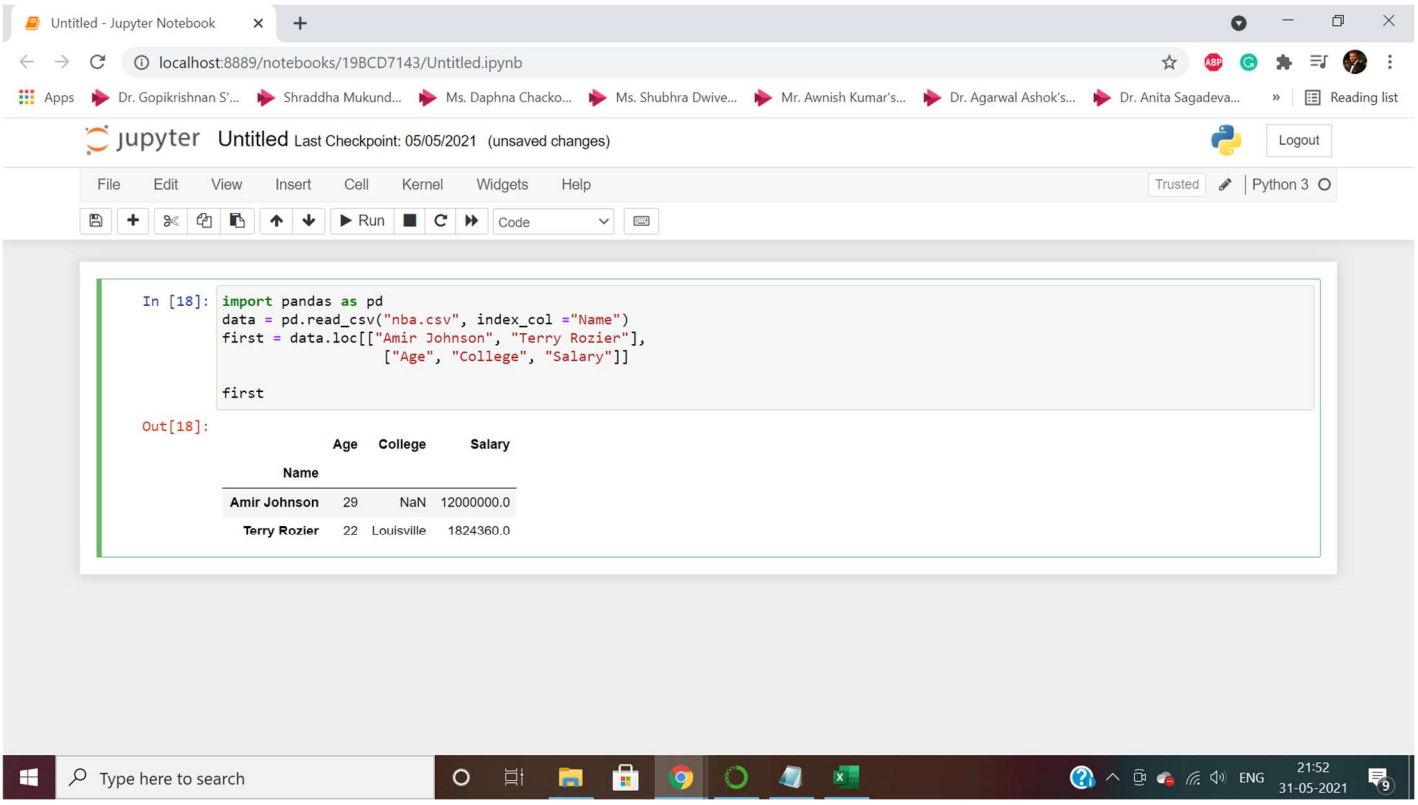

iv) To select row Amir Jhonson, Terry Rozier and John Holland with all columns in a dataframe.

Answer:

```
import pandas as pd
data = pd.read_csv("nba.csv", index_col ="Name")
first = data.loc[["Amir Johnson", "Terry Rozier","John Holland"]]

first
```

OUTPUT

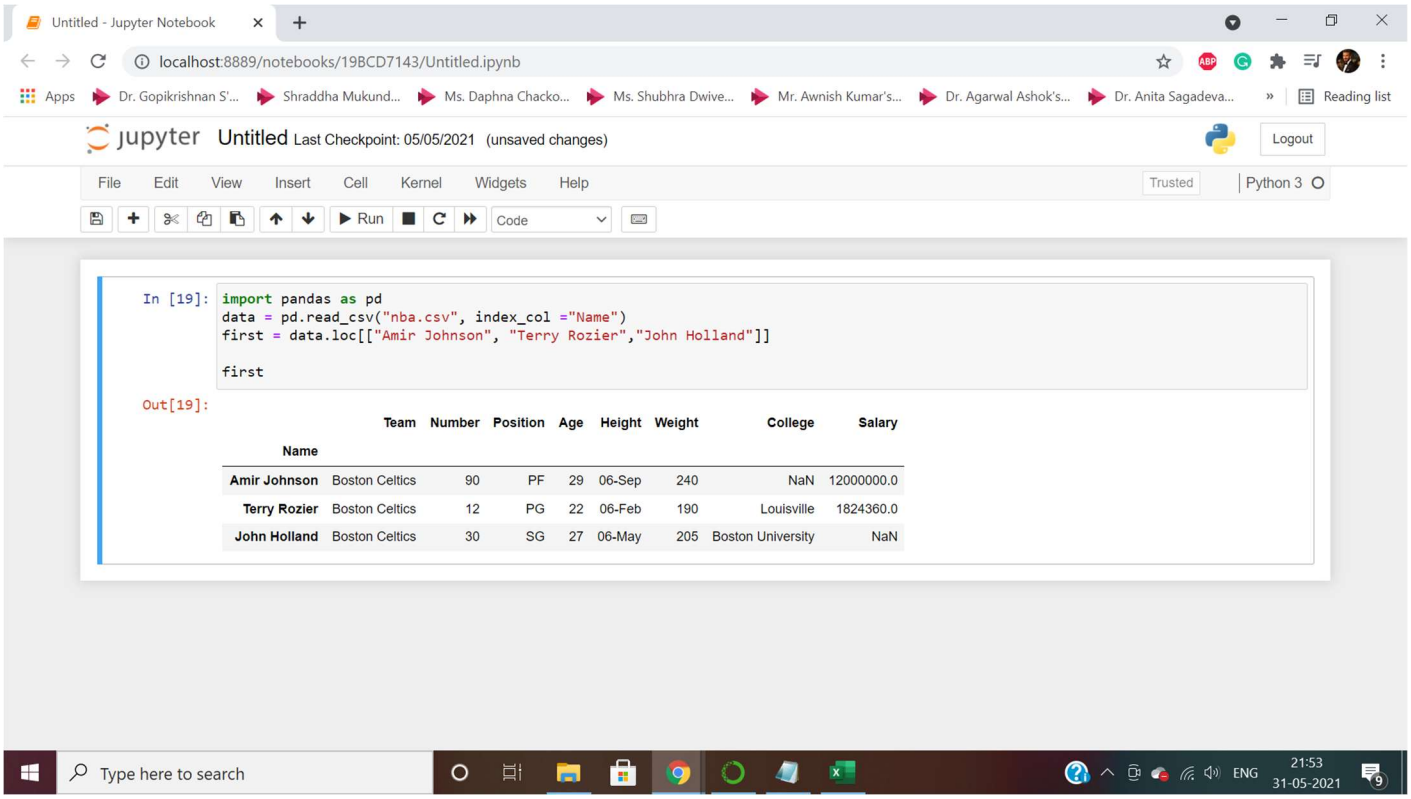

Name	Team	Number	Position	Age	Height	Weight	College	Salary
Amir Johnson	Boston Celtics	90	PF	29	06-Sep	240	NaN	12000000.0
Terry Rozier	Boston Celtics	12	PG	22	06-Feb	190	Louisville	1824360.0
John Holland	Boston Celtics	30	SG	27	06-May	205	Boston University	NaN

v) To select columns Age, Height and Salary with all rows in a dataframe.

Answer:

```
import pandas as pd
data = pd.read_csv("nba.csv", index_col ="Name")
first = data[["Age", "Height", "Salary"]]

first
```

OUTPUT

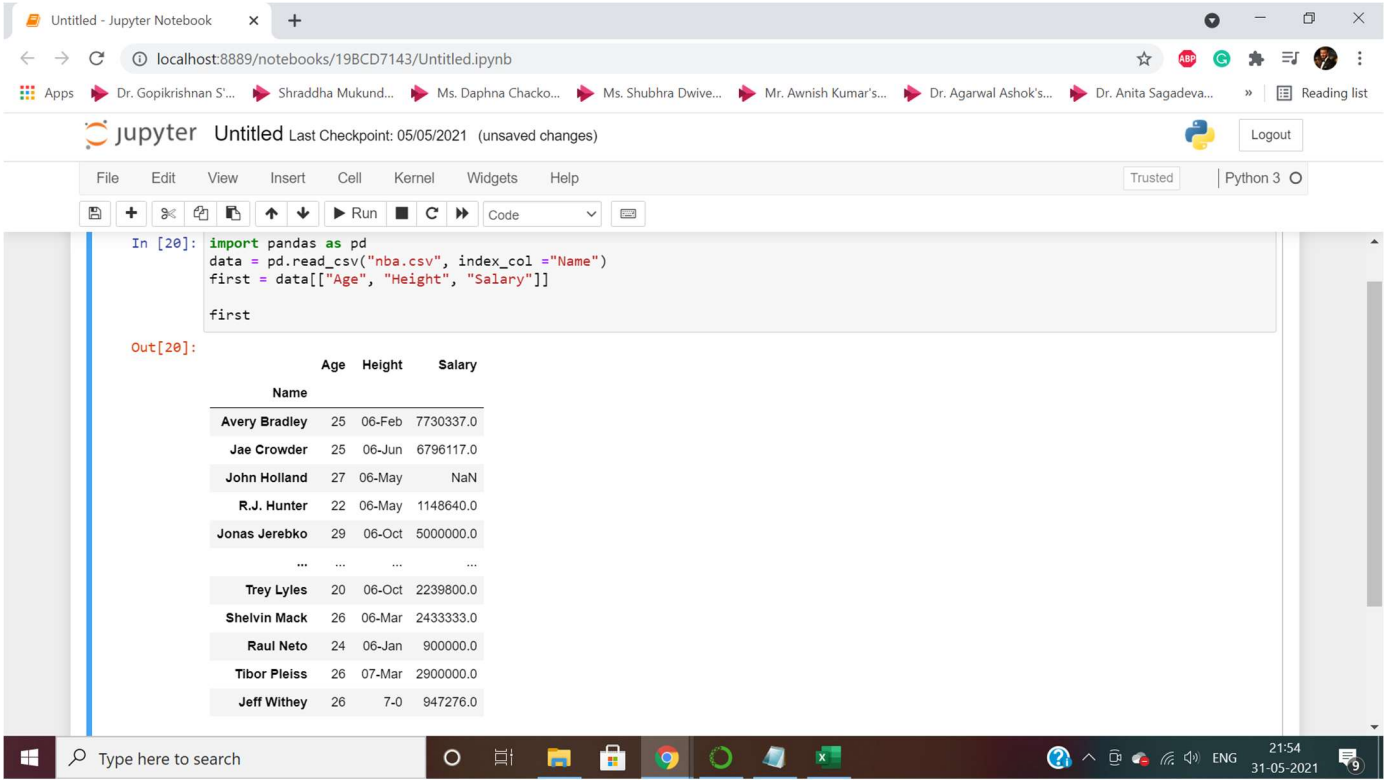

vi) To select column Age, with all rows in a dataframe with label is Name.

Asnwer:

```
import pandas as pd
data = pd.read_csv("nba.csv", index_col ="Name")
first = data["Age"]
first
```

OUTPUT

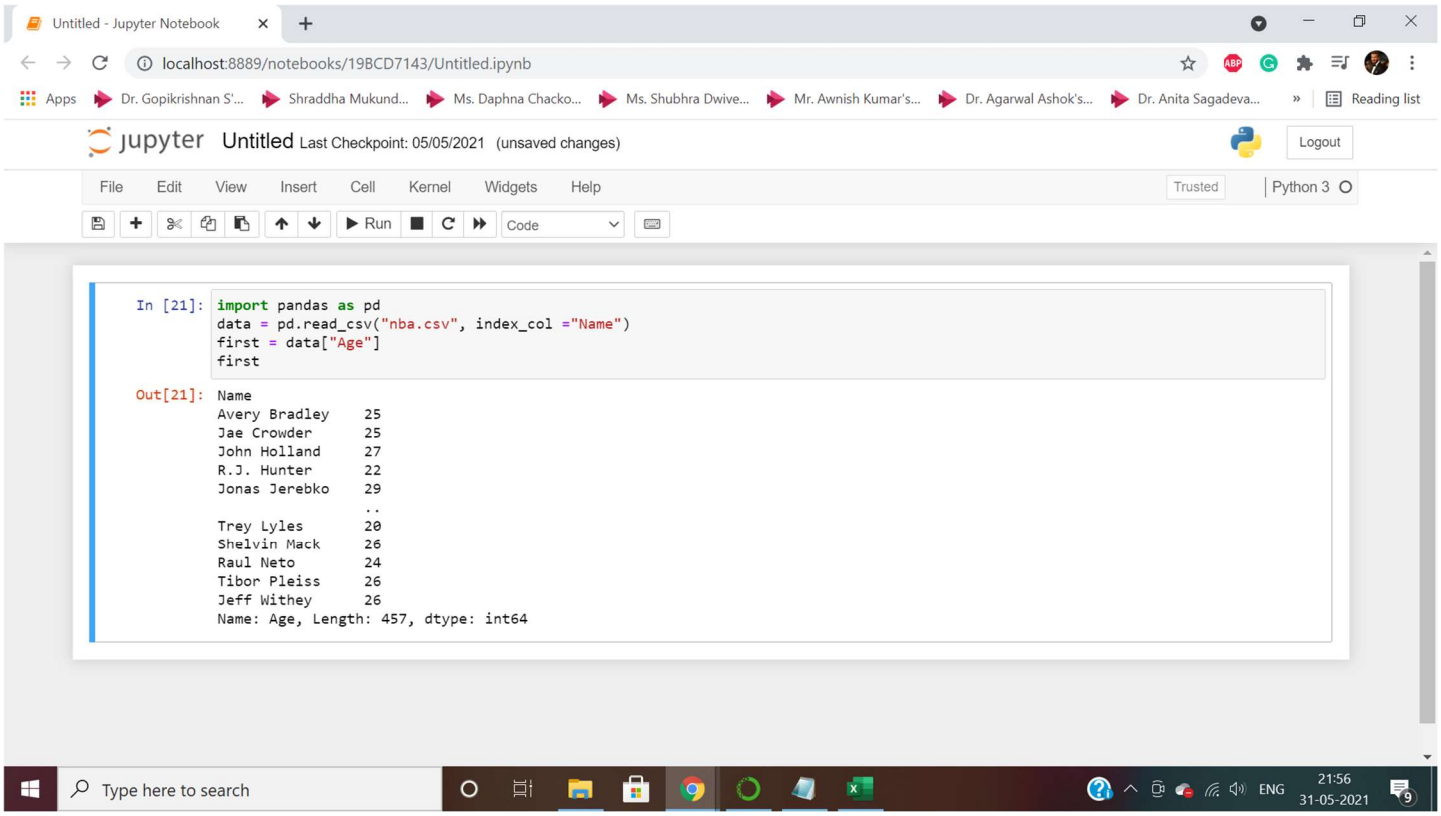

2.Create a dataframe as df = pd.DataFrame(np.random.randn(6, 3), index = ['a','b','c','d','e','f'], columns = ['A', 'B', 'C'])

Find the output of the following syntax:

- **print (df.loc['a':'f'])**
- **print (df.loc['a']>0)**
- **print df.loc[:,'B']**
- **print (df1.iloc[:4])**
- **print (df1.iloc[2:4, 1:3])**

Answer:

import pandas as pd

```
import numpy as np
df =pd.DataFrame(np.random.randn(6, 3), index =
['a','b','c','d','e','f'],columns = ['A', 'B', 'C'])
print ("1)",df.loc['a':'f'])
print("\r")
print ("2)",df.loc['a']>0)
print("\r")
print ("3)",df.loc[:,'B'])
print("\r")
print ("4)",df.iloc[:4])
print("\r")
print ("5)",df.iloc[2:4, 1:3])
```

OUTPUT

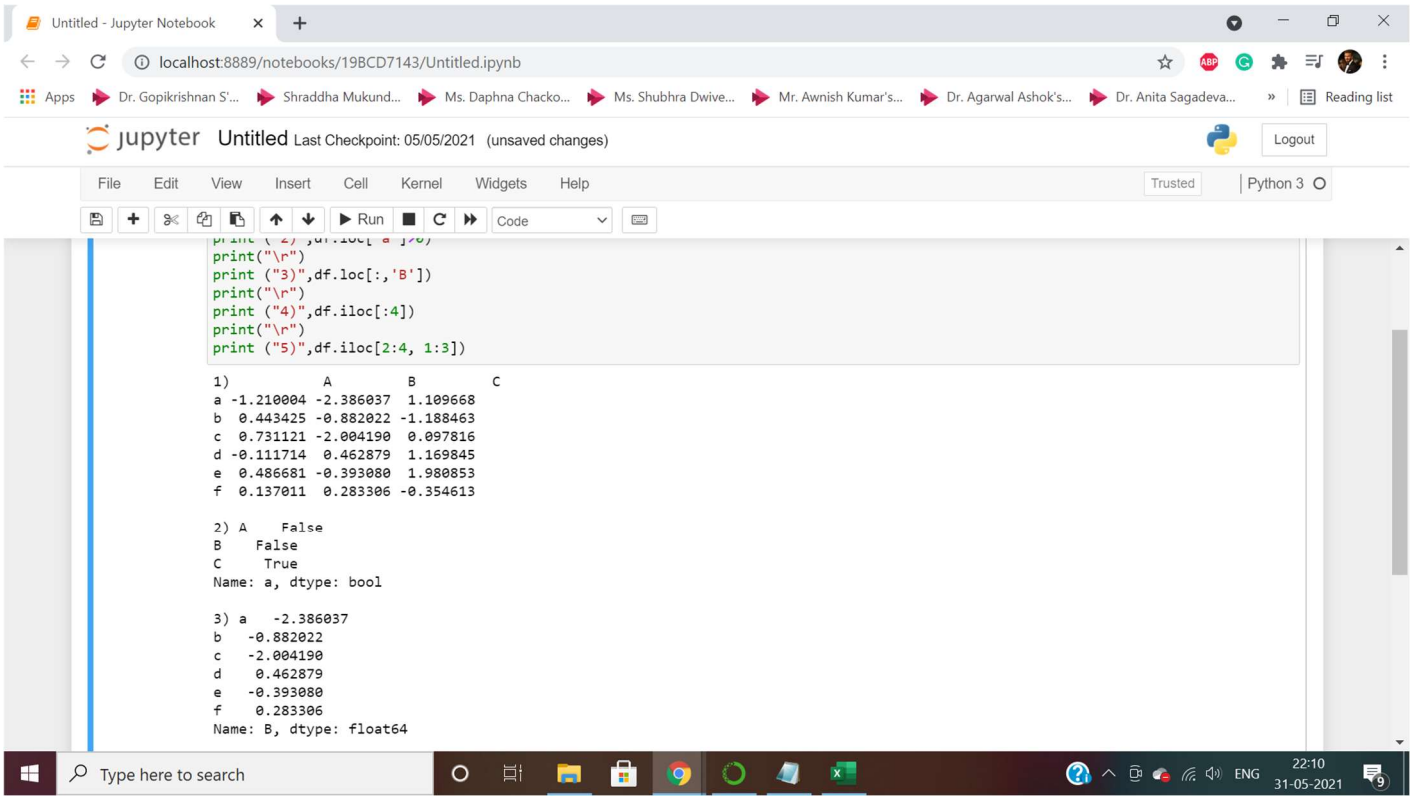

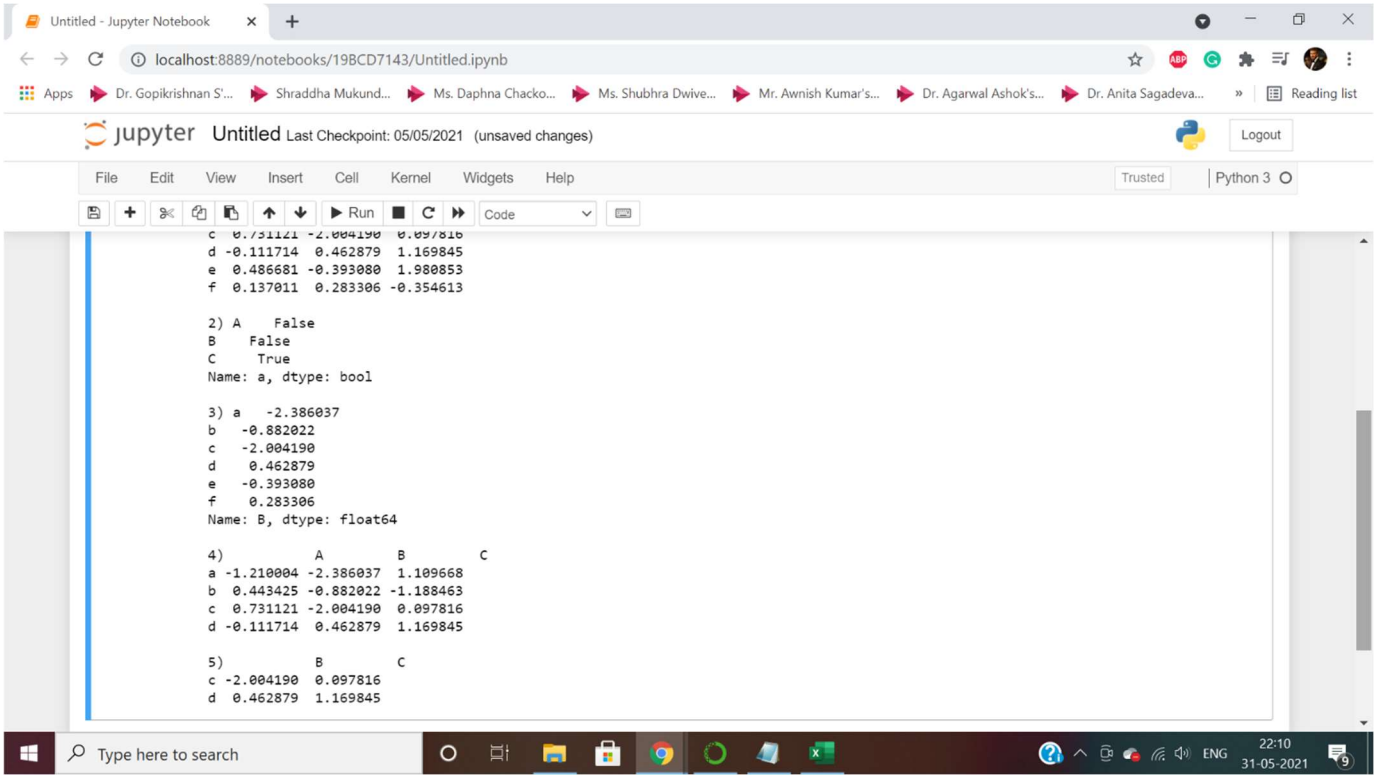

3.Create a dataframe as df5 = pd.DataFrame({'a': ['one', 'one', 'two', 'two', 'two'], 'b': ['x', 'y', 'x', 'y', 'x'], 'c': np.random.randn(5)})

Find the output of the following syntax:

- **Print df5**
- **Print df5.duplicated('a')**
- **Print df5.drop_duplicates('a')**

Answer:

```
import pandas as pd
import numpy as np
df5 = pd.DataFrame({'a': ['one', 'one', 'two', 'two','two'], 'b': ['x', 'y', 'x',
'y', 'x'],'c': np.random.randn(5)})
print ("1)",df5)
print("\r")
print ("2)",df5.duplicated('a'))
print("\r")
```

print ("3)",df5.drop_duplicates('a'))

OUTPUT

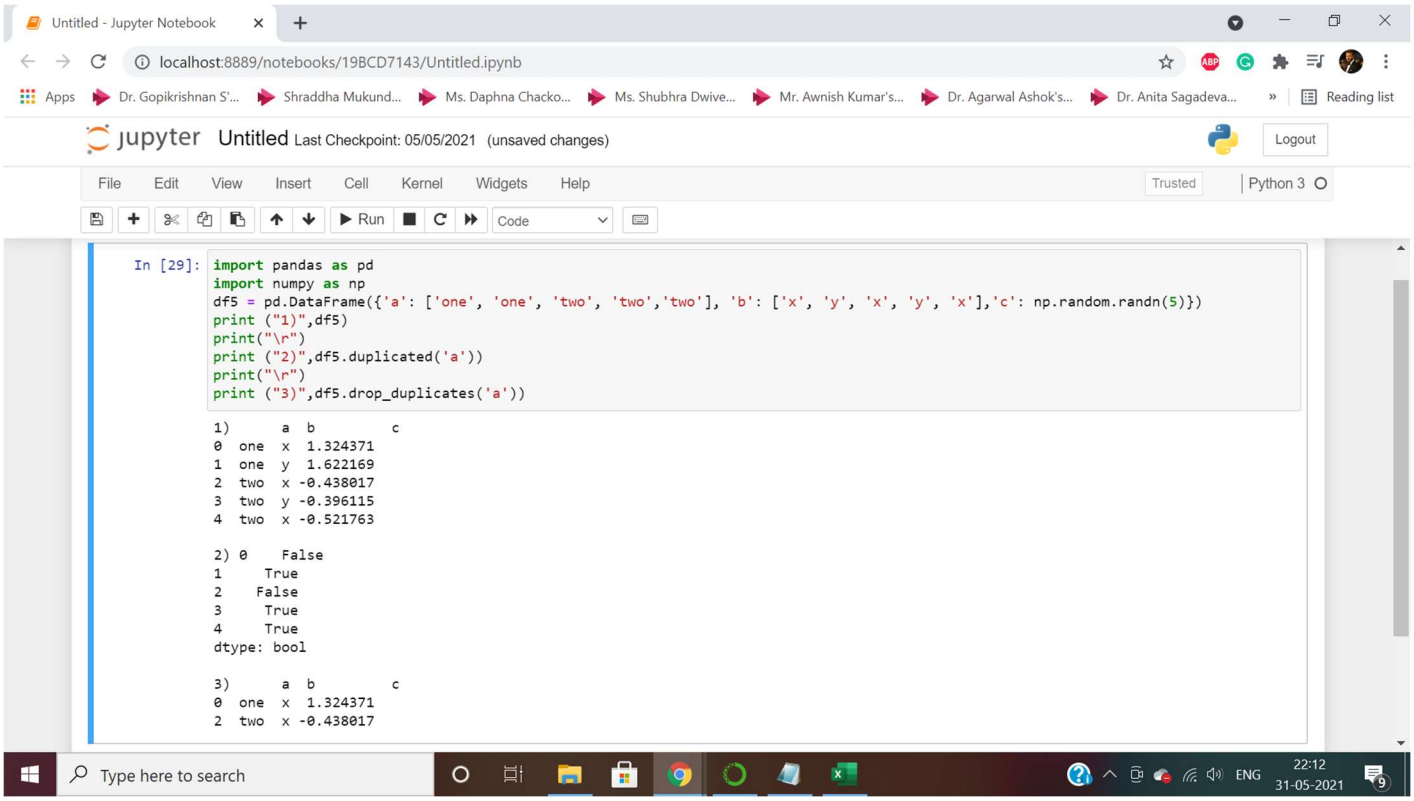

4. Create the following series data = pd.Series(['a', 'b', 'c'], index=[1, 3, 5])

Find the output of the following syntax:

· Print data · Print data[1]

· Print data[1:3]

· Print data.loc[1]

· Print data.loc[1:3]

· Print data.iloc[1:3]

Answer:

```
import pandas as pd
import numpy as np
data = pd.Series(['a', 'b', 'c'], index=[1, 3, 5])
print ("1)",data)
print("\r")
print ("2)",data[1])
print("\r")
print ("3)",data[1:3])
```

```
print("\r")
print ("4)",data.loc[1])
print("\r")
print ("5)",data.loc[1:3])
print("\r")
print ("6)",data.iloc[1:3])
```

OUTPUT

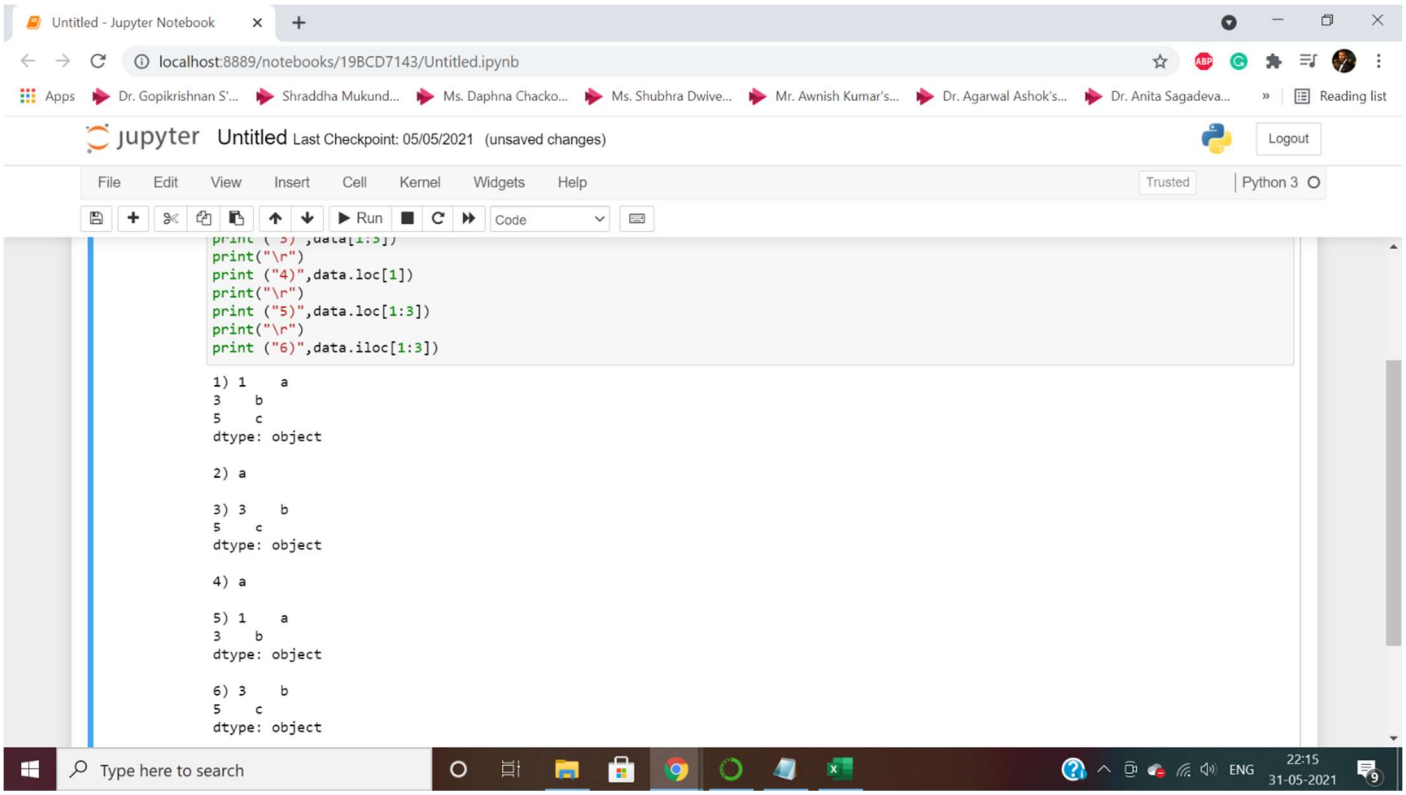

5. Create the following series area = pd.Series({'California': 423967, 'Texas': 695662, 'New York': 141297, 'Florida': 170312, 'Illinois': 149995}) pop = pd.Series({'California': 38332521, 'Texas': 26448193, 'New York': 19651127, 'Florida': 19552860, 'Illinois': 12882135})

Find the output of the following syntax:

· Create dataframe by following syntax as: data = pd.DataFrame({'area':area, 'pop':pop})

· Print data · Print data['area']

· Print data['density'] = data['pop'] / data['area']

Answer:

```
import pandas as pd
import numpy as np
area = pd.Series({'California': 423967, 'Texas': 695662,
'New York': 141297, 'Florida': 170312,
'Illinois': 149995})
pop = pd.Series({'California': 38332521, 'Texas': 26448193,
'New York': 19651127, 'Florida': 19552860,
'Illinois': 12882135})

data =pd.DataFrame({'area':area, 'pop':pop})

print ("1)",data)
print("\r")
print ("2)",data['area'])
print("\r")
data['density'] = data['pop'] / data['area']
print ("3)",data)
```

OUTPUT

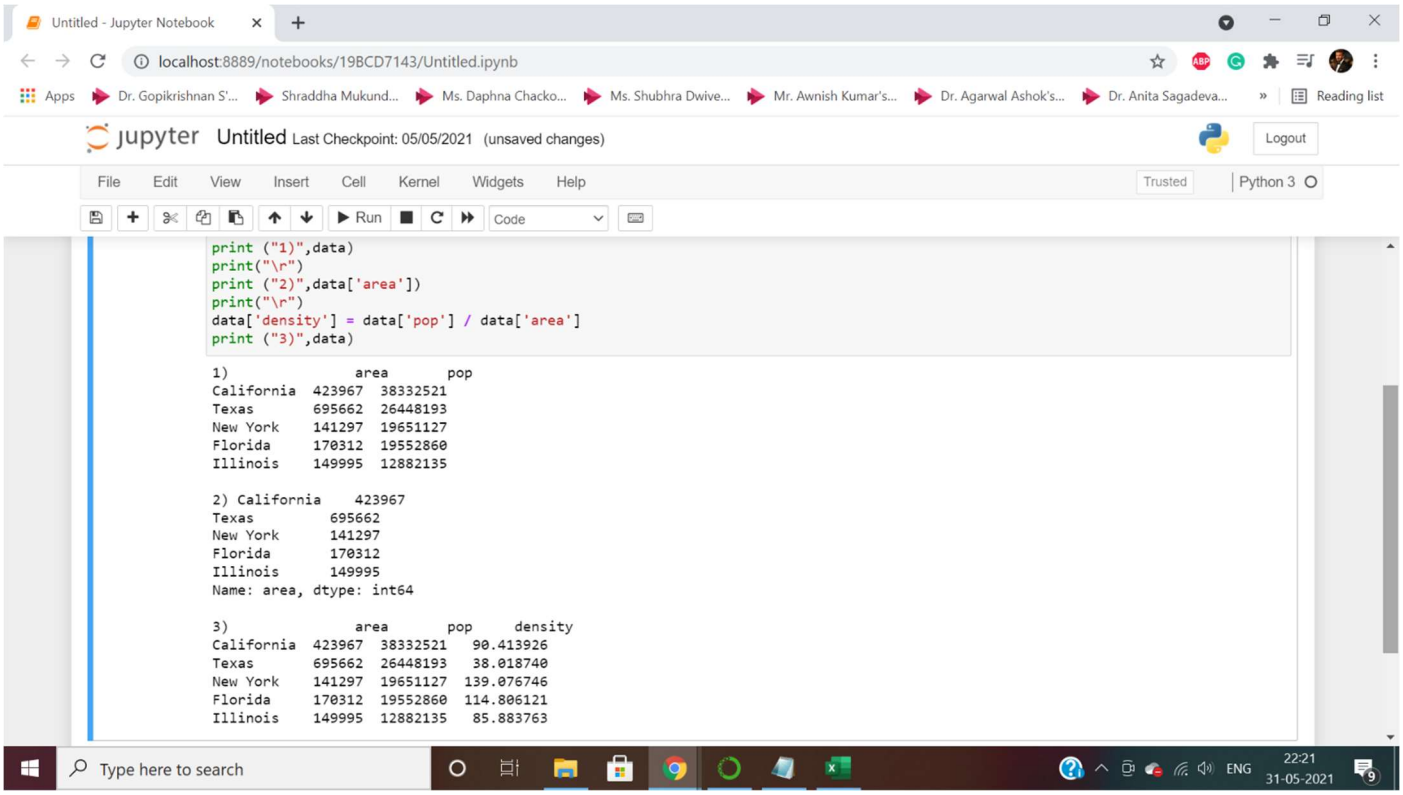

6. . Create the following series Create as index = [('California', 2000), ('California', 2010), ('New York', 2000), ('New York', 2010), ('Texas', 2000), ('Texas', 2010)] populations = [33871648, 37253956, 18976457, 19378102, 20851820, 25145561] pop = pd.Series(populations, index=index)

Find the output of the following syntax:

· Print pop

· Print pop[('California', 2010):('Texas', 2000)]

· Print pop[[i for i in pop.index if i[1] == 2010]]

· Print index = pd.MultiIndex.from_tuples(index)

Answer:

```
import pandas as pd
import numpy as np
index = [('California', 2000), ('California', 2010),
('New York', 2000), ('New York', 2010),
('Texas', 2000), ('Texas', 2010)]
populations = [33871648, 37253956,
18976457, 19378102,
20851820, 25145561]
pop = pd.Series(populations, index=index)

print ("1)",pop)
print("\r")
print ("2)",pop[('California', 2010):('Texas', 2000)])
print("\r")
print ("3)",pop[[i for i in pop.index if i[1] == 2010]])
print("\r")
index = pd.MultiIndex.from_tuples(index)
print ("4)",index)
```

OUTPUT

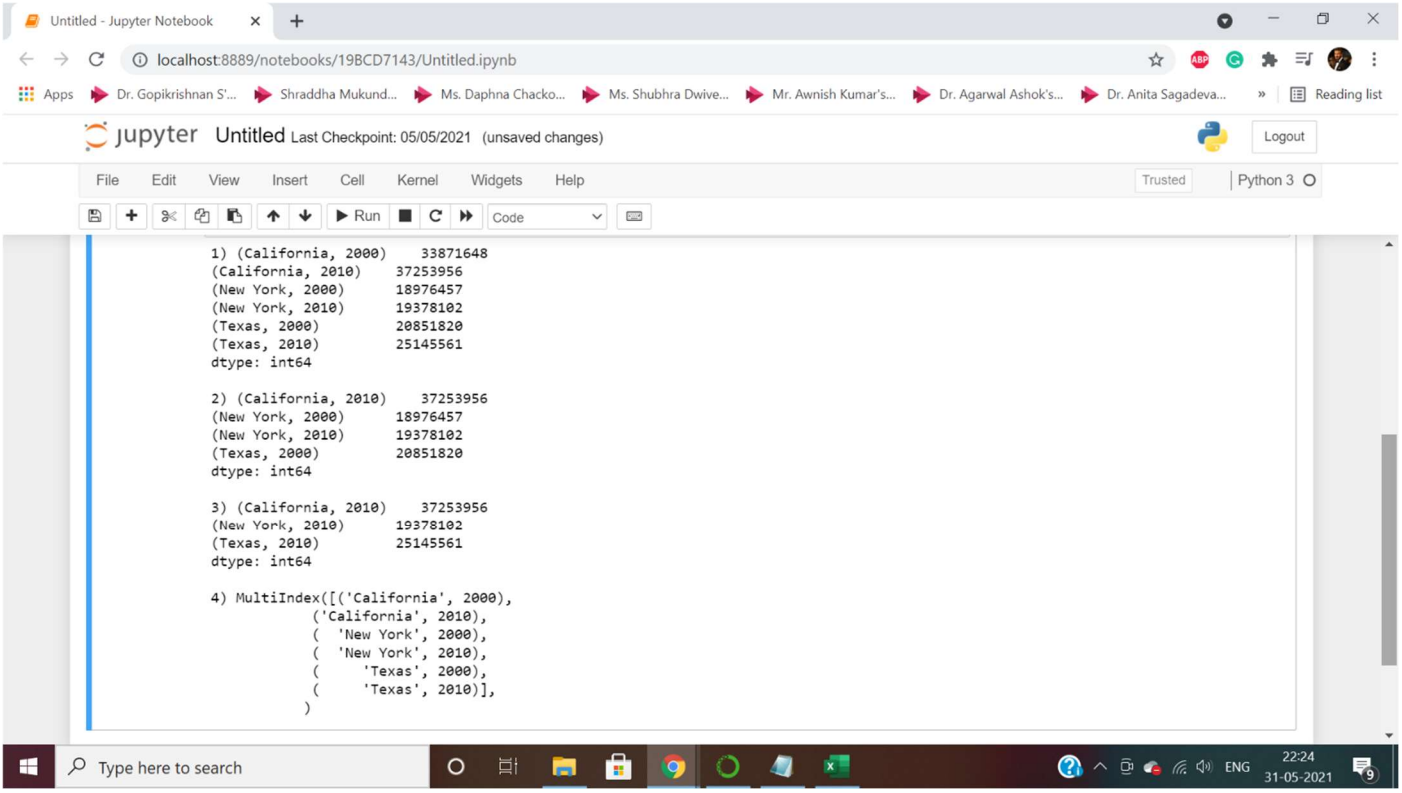

7. Create a dataframe as df = pd.DataFrame([[1, np.nan, 2], [2, 3, 5], [np.nan, 4, 6]])

Find the output of the following syntax:

· Print df · Print df.dropna()

· Print df.dropna(axis='columns')

· Print df.dropna(axis='columns', how='all')

Answer:

```
import pandas as pd
import numpy as np
df = pd.DataFrame([[1, np.nan, 2],
[2, 3, 5],
[np.nan, 4, 6]])

print ("1)",df)
```

```
print("\r")
print ("2)",df.dropna())
print("\r")
print ("3)",df.dropna(axis='columns'))
print("\r")
print ("4)",df.dropna(axis='columns', how='all'))
```

OUTPUT

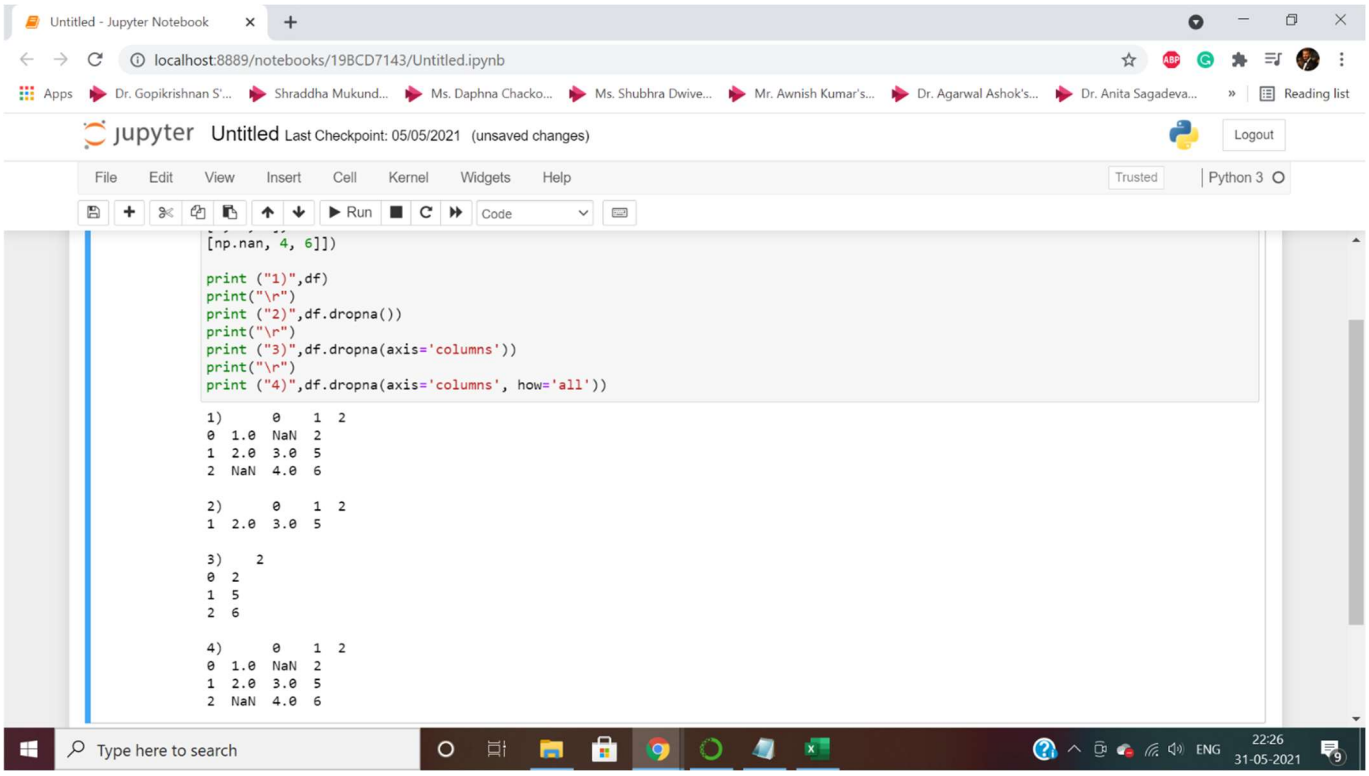

8.

a. Consider the following dataframe : student_df and write a statement of below mentioned frame to get the minimum value of the column marks:

Name Course Marks
Anamay FDA 95
Aditi FDA 82
Mehak FDA 65
Kriti FDA 55

Answer:

```
import pandas as pd
```

```
import numpy as np

data = {
  "Name": ["Anamay", "Aditi", "Mehak","Kriti"],
  "Course": ["FDA", "FDA", "FDA", "FDA"],
  "Marks": [95, 82, 65, 55]}

student_df = pd.DataFrame(data)

print (student_df)
print("\r")
print ("Minimum value of the column marks:",student_df.min().Marks)
```

OUTPUT

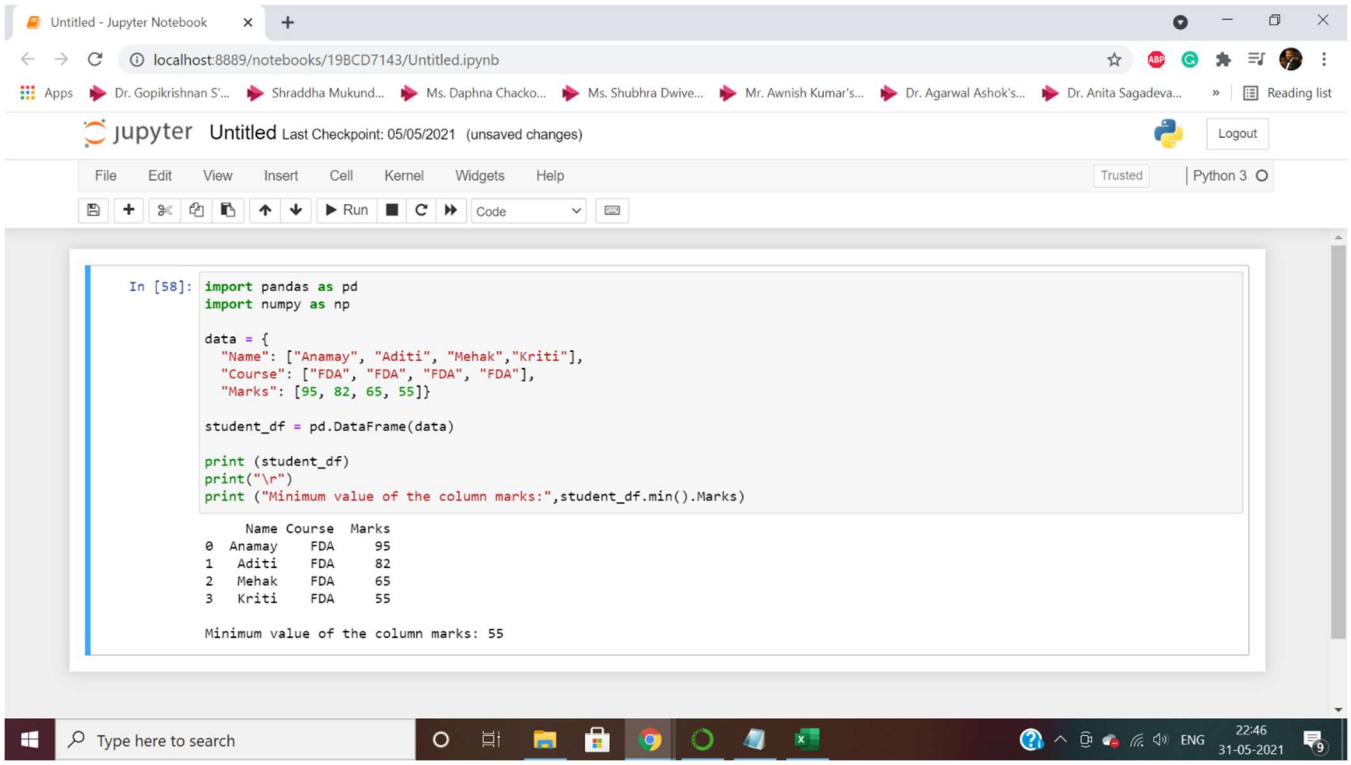

b. Write the output of the following syntax:

```
import numpy as np
array1=np.array([10,12,14,16,18,20,22])
print(array1[1:5:2])
```

Answer:

```
import pandas as pd
import numpy as np

array1=np.array([10,12,14,16,18,20,22])
print(array1[1:5:2])
```

OUTPUT

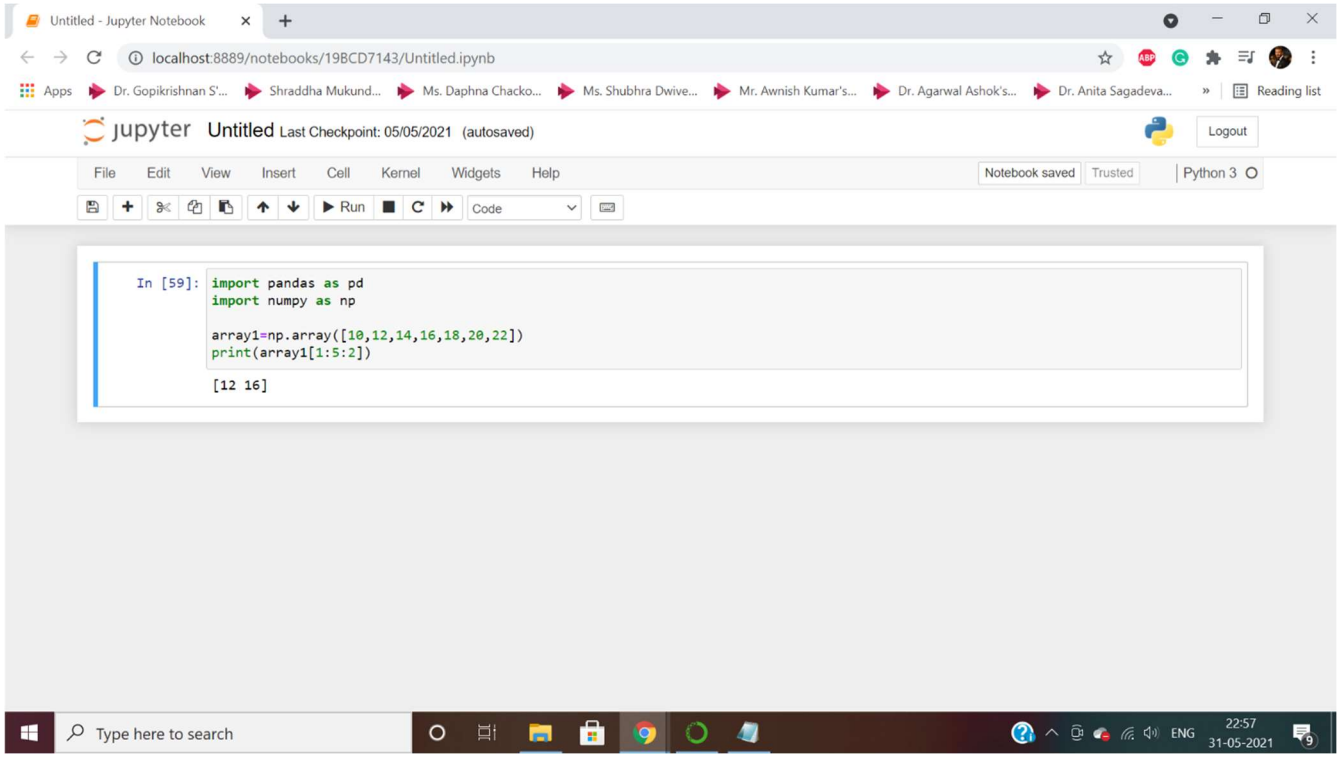

c. Consider the following dataframe, and answer the questions given below: import pandas as pd df = pd.DataFrame({"Quarter1":[2000, 4000, 5000, 4400, 10000], "Quarter2":[5800, 2500, 5400, 3000, 2900], "Quarter3":[20000, 16000, 7000, 3600, 8200], "Quarter4":[1400, 3700, 1700, 2000, 6000]})

(i) Write the code to find mean value from above dataframe df over the index and column axis.

(ii) Use sum() function to find the sum of all the values over the index axis.

Answer:

```
import pandas as pd
import numpy as np
```

```
df = pd.DataFrame({"Quarter1":[2000, 4000, 5000, 4400, 10000],
"Quarter2":[5800, 2500, 5400, 3000, 2900], "Quarter3":[20000, 16000,
7000,
3600, 8200], "Quarter4":[1400, 3700, 1700, 2000, 6000]})

print("i)")
print("mean value from above dataframe df over the index")
print(df.mean(axis = 1, skipna = True))
print("\r")
print("mean value from above dataframe df over the column axis")
print(df.mean(axis = 0, skipna = True))
print("\r")
print("ii)")
print("Sum of all the values over the index axis")
print(df.sum(axis = 1, skipna = True))
```

OUTPUT

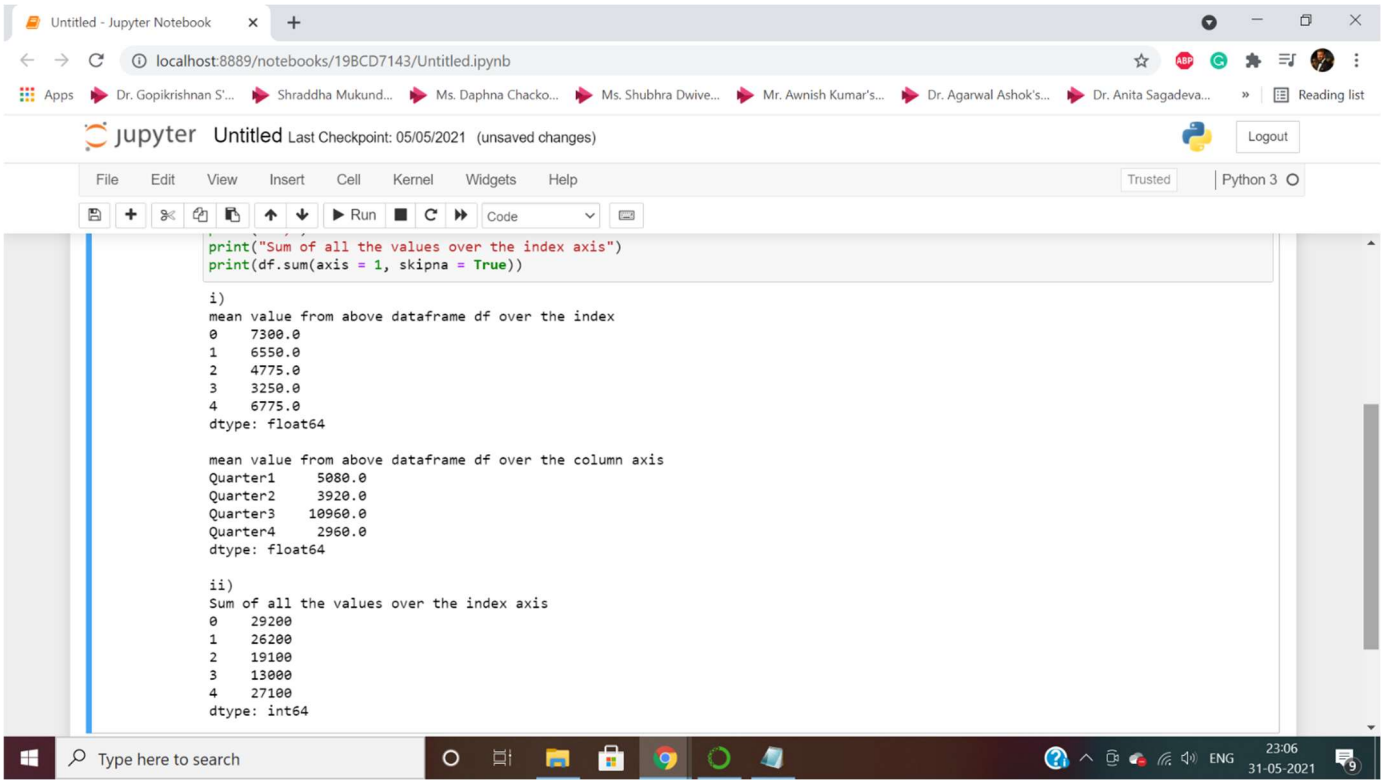

9. In this question (9), we will use attached dataset as pima-indians-diabetes.csv.

· load the dataset
· summarize the dataset
· print the first 20 rows of data
· count the number of missing values for each column
· fill missing values with mean column values
· replace '0' values with 'nan
· count the number of nan values in each column
· drop rows with missing values

Answer:

1)import pandas as pd
data = pd.read_csv("pima-indians-diabetes.csv")
data.head()

OUTPUT

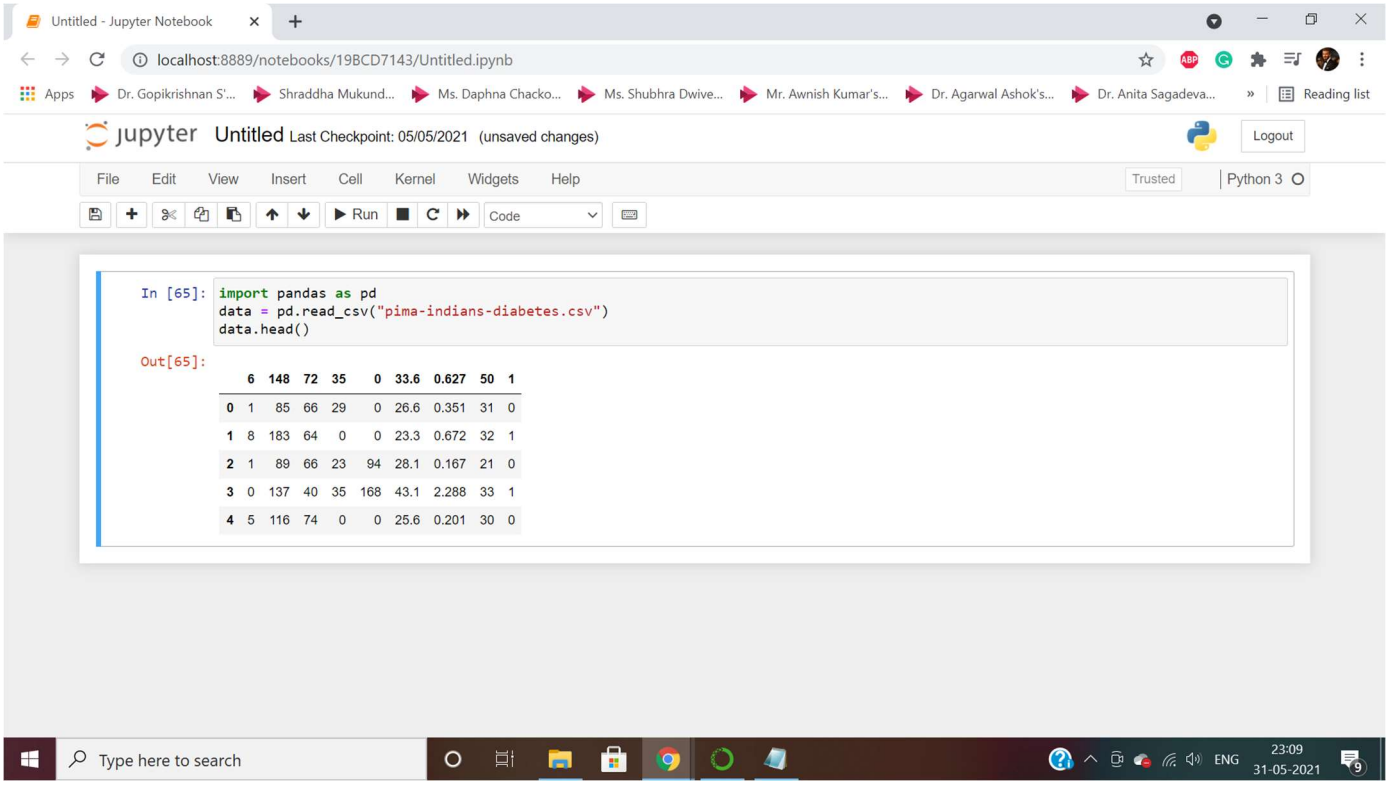

	6	148	72	35	0	33.6	0.627	50	1
0	1	85	66	29	0	26.6	0.351	31	0
1	8	183	64	0	0	23.3	0.672	32	1
2	1	89	66	23	94	28.1	0.167	21	0
3	0	137	40	35	168	43.1	2.288	33	1
4	5	116	74	0	0	25.6	0.201	30	0

2)import pandas as pd
data = pd.read_csv("pima-indians-diabetes.csv")
data.describe()

OUTPUT

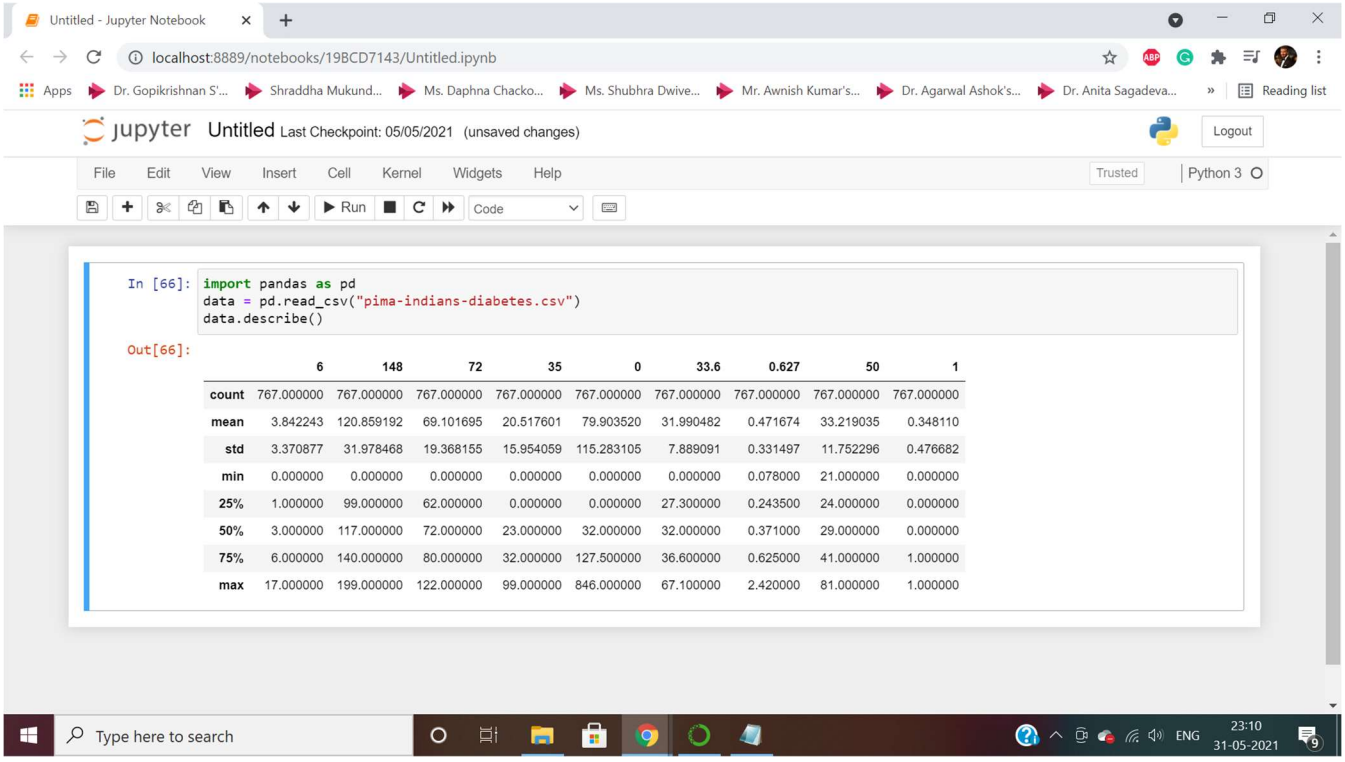

	6	148	72	35	0	33.6	0.627	50	1
count	767.000000	767.000000	767.000000	767.000000	767.000000	767.000000	767.000000	767.000000	767.000000
mean	3.842243	120.859192	69.101695	20.517601	79.903520	31.990482	0.471674	33.219035	0.348110
std	3.370877	31.978468	19.368155	15.954059	115.283105	7.889091	0.331497	11.752296	0.476682
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243500	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	32.000000	32.000000	0.371000	29.000000	0.000000
75%	6.000000	140.000000	80.000000	32.000000	127.500000	36.600000	0.625000	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

3)import pandas as pd

data = pd.read_csv("pima-indians-diabetes.csv")
data.head(20)

OUTPUT

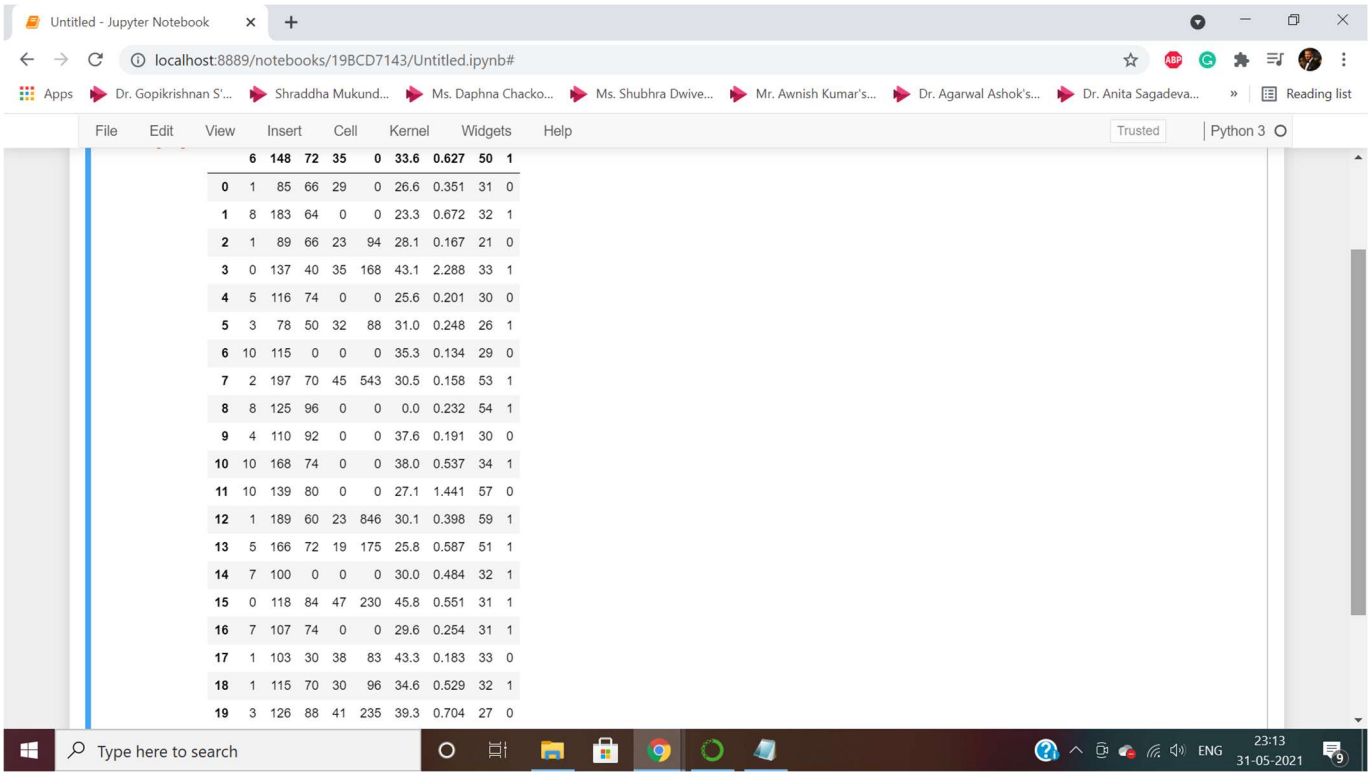

	6	148	72	35	0	33.6	0.627	50	1
0	1	85	66	29	0	26.6	0.351	31	0
1	8	183	64	0	0	23.3	0.672	32	1
2	1	89	66	23	94	28.1	0.167	21	0
3	0	137	40	35	168	43.1	2.288	33	1
4	5	116	74	0	0	25.6	0.201	30	0
5	3	78	50	32	88	31.0	0.248	26	1
6	10	115	0	0	0	35.3	0.134	29	0
7	2	197	70	45	543	30.5	0.158	53	1
8	8	125	96	0	0	0.0	0.232	54	1
9	4	110	92	0	0	37.6	0.191	30	0
10	10	168	74	0	0	38.0	0.537	34	1
11	10	139	80	0	0	27.1	1.441	57	0
12	1	189	60	23	846	30.1	0.398	59	1
13	5	166	72	19	175	25.8	0.587	51	1
14	7	100	0	0	0	30.0	0.484	32	1
15	0	118	84	47	230	45.8	0.551	31	1
16	7	107	74	0	0	29.6	0.254	31	1
17	1	103	30	38	83	43.3	0.183	33	0
18	1	115	70	30	96	34.6	0.529	32	1
19	3	126	88	41	235	39.3	0.704	27	0

4)import pandas as pd
data = pd.read_csv("pima-indians-diabetes.csv")
print("Number of missing values of the said dataframe:")
print(data.isna().sum())

OUTPUT

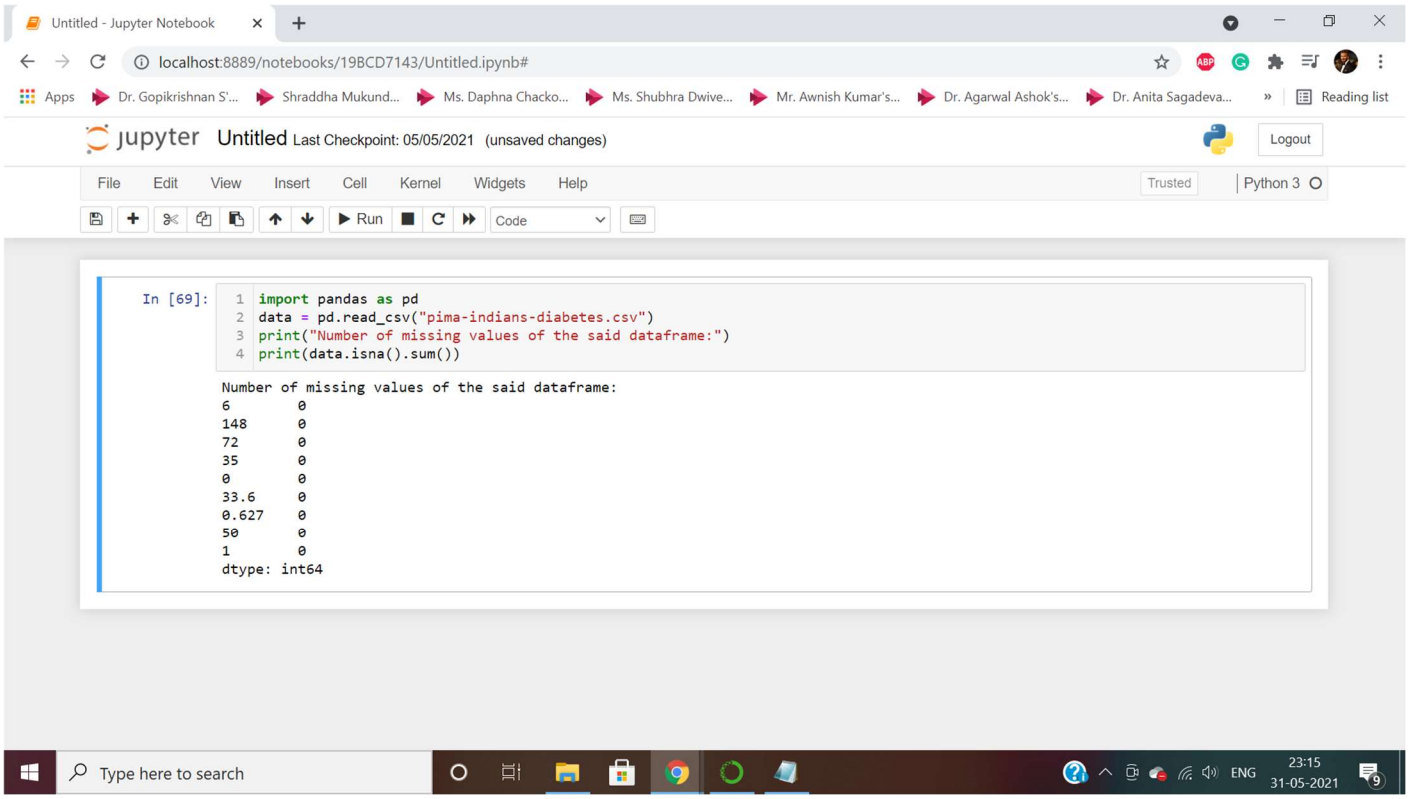

5)There are no missing values,so we cant fill it with mean values

6)import pandas as pd

data = pd.read_csv("pima-indians-diabetes.csv")

data[["6","148","72","35","0","33.6","0.627","50","1"]].astype(str).replace('0', np.nan)

OUTPUT

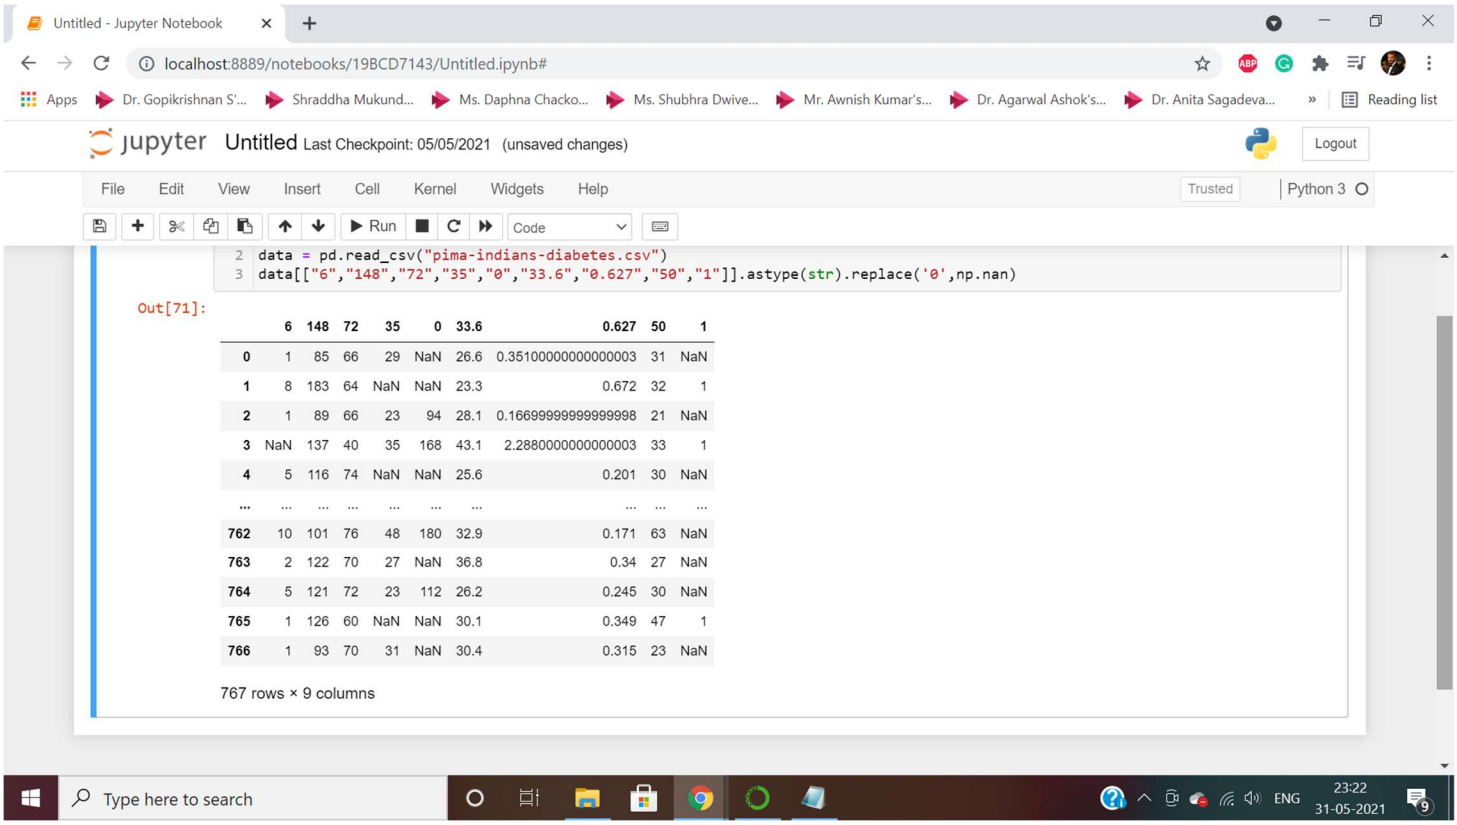

	6	148	72	35	0	33.6	0.627	50	1
0	1	85	66	29	NaN	26.6	0.35100000000000003	31	NaN
1	8	183	64	NaN	NaN	23.3	0.672	32	1
2	1	89	66	23	94	28.1	0.16699999999999998	21	NaN
3	NaN	137	40	35	168	43.1	2.2880000000000003	33	1
4	5	116	74	NaN	NaN	25.6	0.201	30	NaN
...	...	...	...	...	...	...	...	...	...
762	10	101	76	48	180	32.9	0.171	63	NaN
763	2	122	70	27	NaN	36.8	0.34	27	NaN
764	5	121	72	23	112	26.2	0.245	30	NaN
765	1	126	60	NaN	NaN	30.1	0.349	47	1
766	1	93	70	31	NaN	30.4	0.315	23	NaN

7)import pandas as pd

```
data = pd.read_csv("pima-indians-diabetes.csv")
data=data[["6","148","72","35","0","33.6","0.627","50","1"]].astype(str).
replace('0',np.nan)
data.isna().sum()
```

OUTPUT

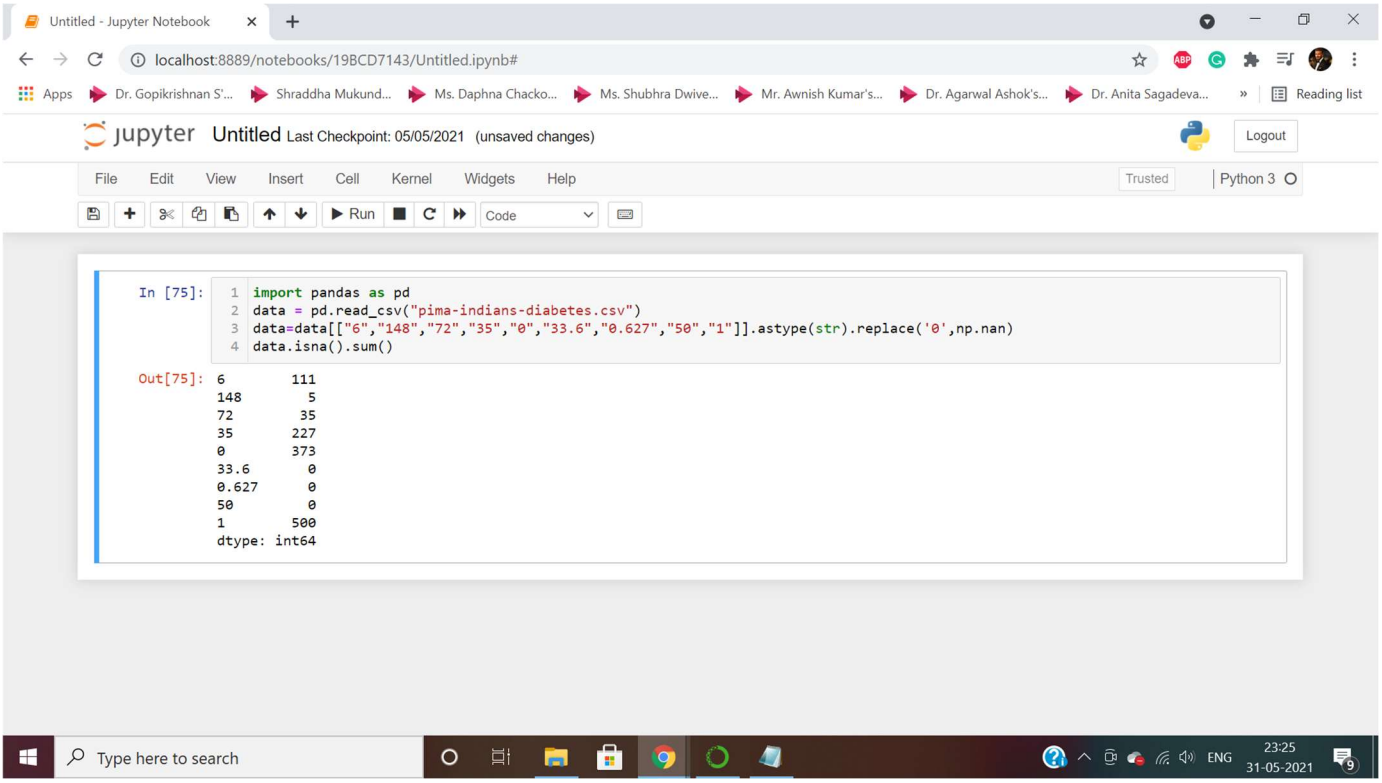

8)import pandas as pd data = pd.read_csv("pima-indians-diabetes.csv")

data=data[["6","148","72","35","0","33.6","0.627","50","1"]].astype(str).replace('0',np.nan)
data.dropna()

OUTPUT

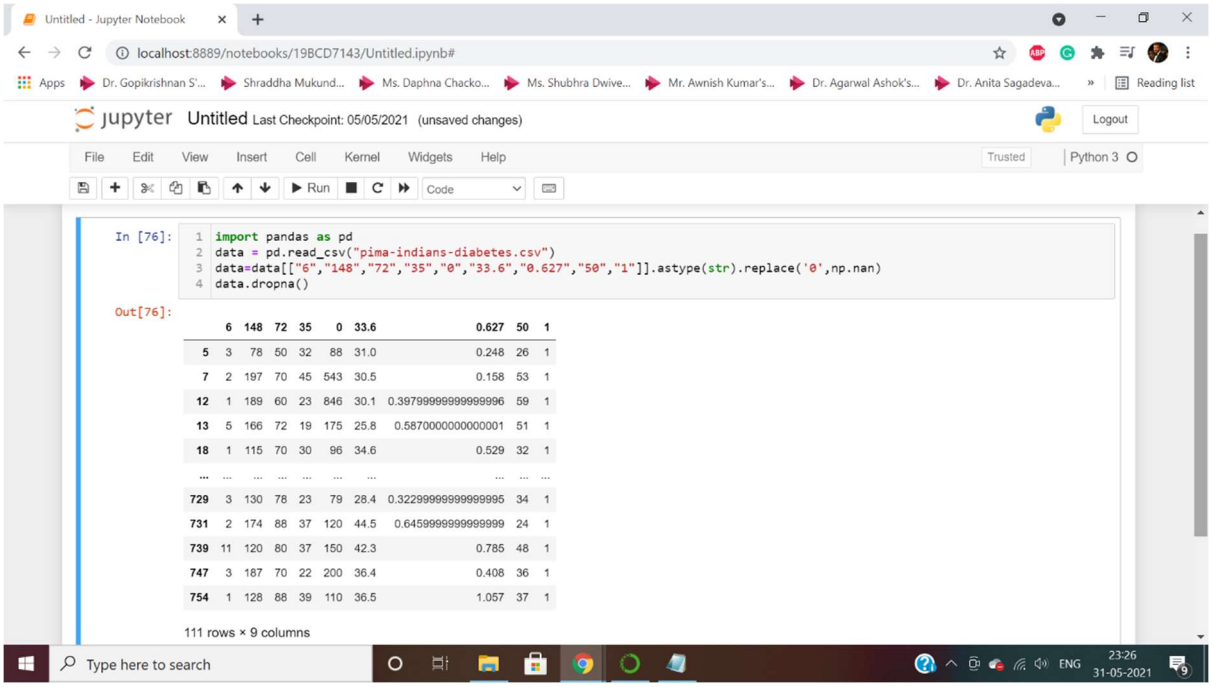

10. Consider the following dataframe, and answer the questions given below:

df = pd.DataFrame(np.random.randn(5, 3), index=['a', 'c', 'e', 'f', 'h'], columns=['one', 'two', 'three']) df = df.reindex(['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'])

· Count the number of nan values in each column

· Replace nan values with 22

Answer:

1)import pandas as pd

df = pd.DataFrame(np.random.randn(5, 3), index=['a', 'c', 'e', 'f', 'h'], columns=['one', 'two', 'three'])

df = df.reindex(['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'])

df.isna().sum()

OUTPUT

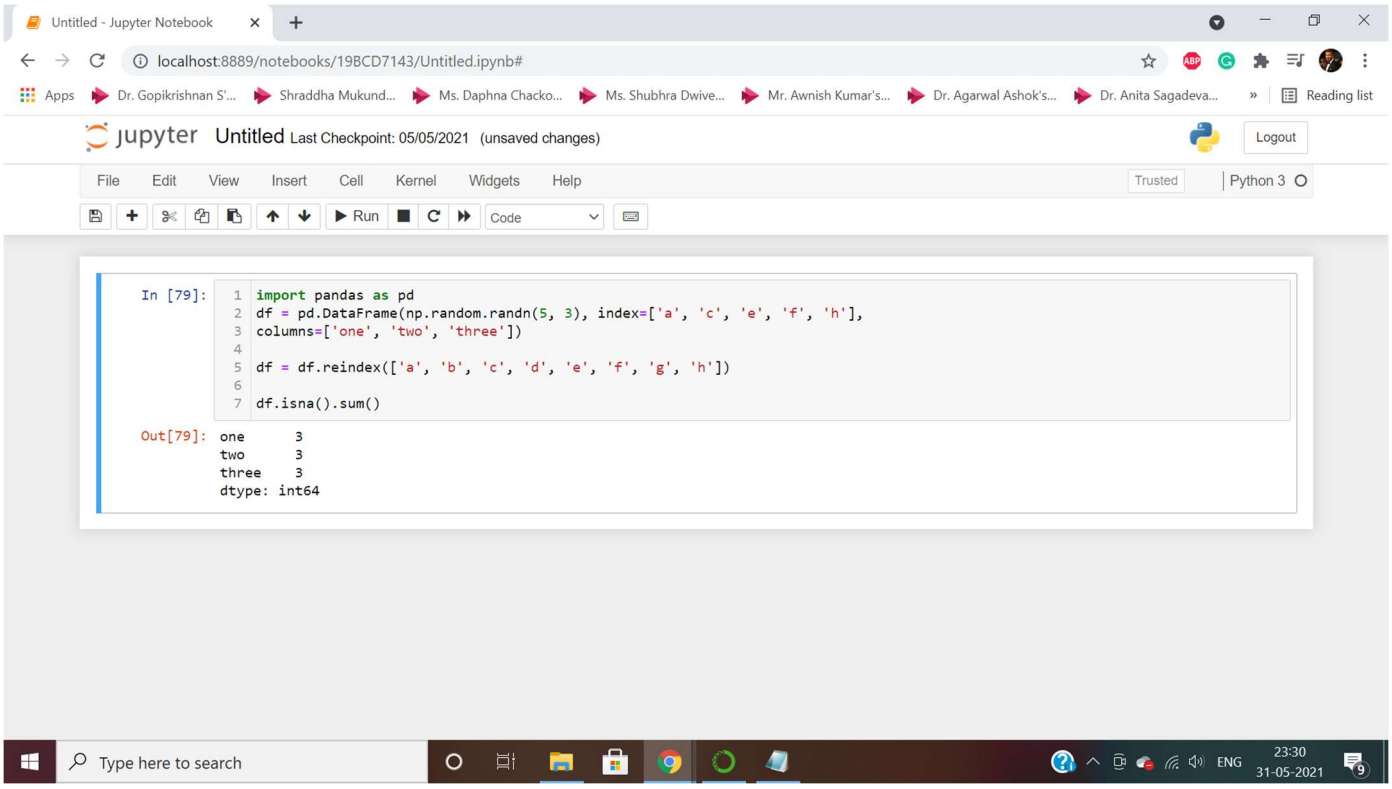

2)import pandas as pd
df = pd.DataFrame(np.random.randn(5, 3), index=['a', 'c', 'e', 'f', 'h'],
columns=['one', 'two', 'three'])

df = df.reindex(['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h'])
df.replace(np.nan,22)

OUTPUT